Finetuned Language Models Are Zero-Shot Learners[1]

September 17, 2021

Abstract

This paper aims to improve the zero-shot learning abilities of language models and proposes **instruction tuning**–finetuning language models on a collection of tasks described via instructions.

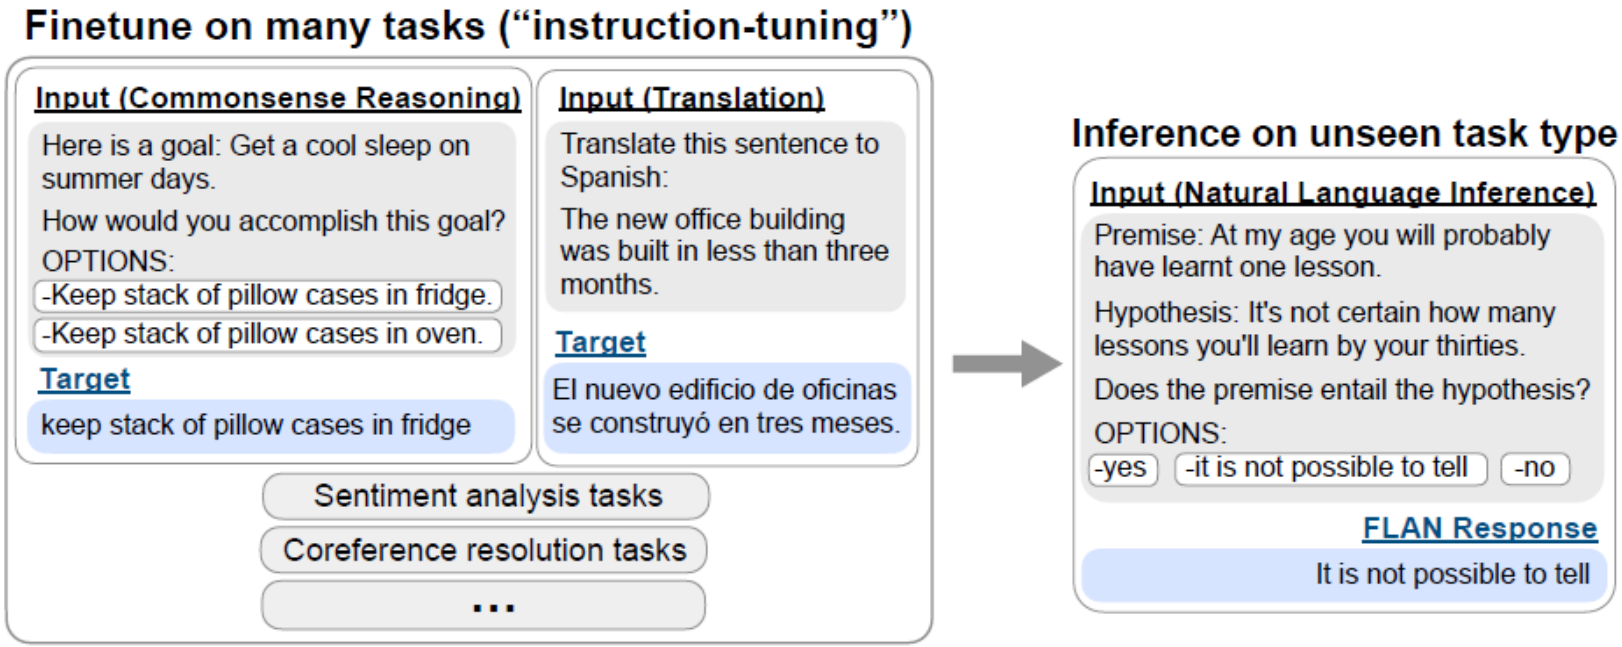

Figure 1: Framework

Overview

GPT-3

GPT-3[2] has been shown to perform few-shot learning remarkably well but less successful at zero-shot learning.

Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.

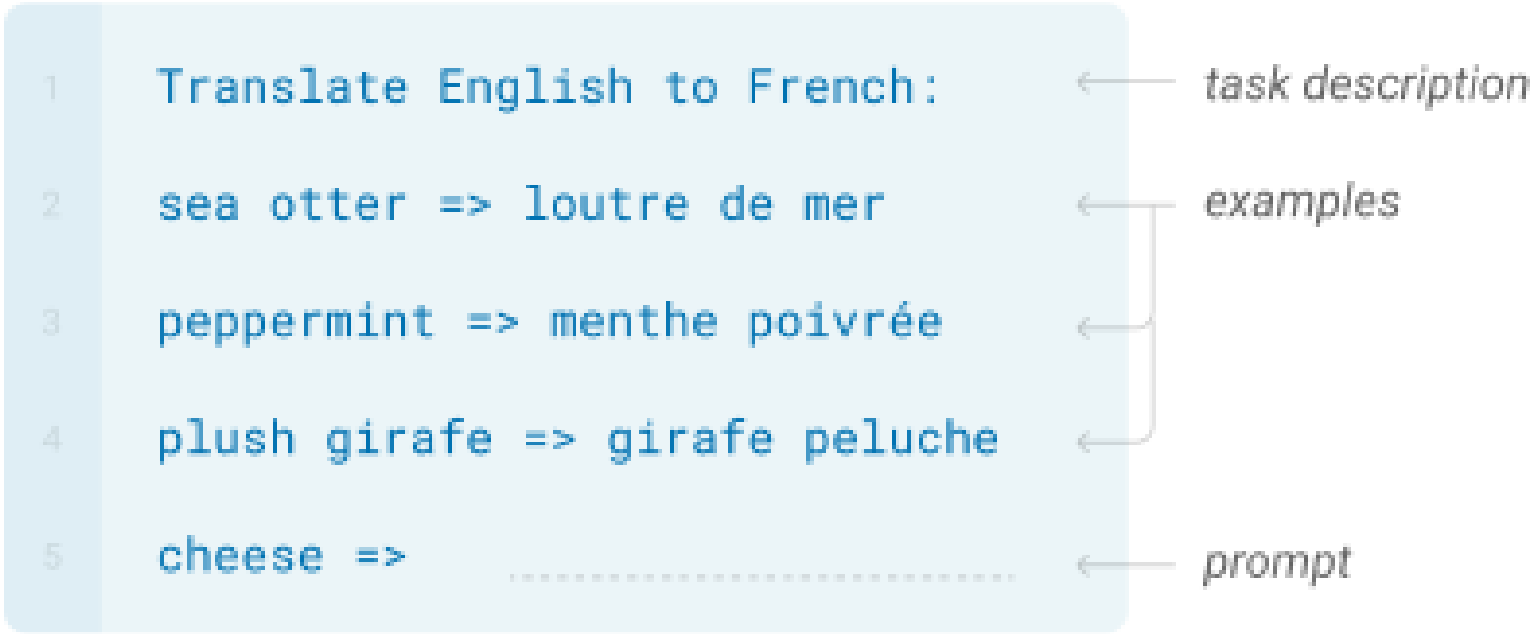

FLAN

The paper leverages the intuition that NLP tasks can be described via natural language instructions, such as:

- "Is the sentiment of this movie review positive or negative?"
- "Translate 'how are you' into Chinese."

The process is as follows:

- pretraining a language model of 137B parameters
- grouping NLP tasks into clusters based on their task types and hold out one for evaluation;
- instruction tuning the LM on all other clusters; the resulting model is called Finetuned LAnguage Net(FLAN).

FLAN

Prompt

```
Explain the moon landing to a 6 year old in a few sentences.
```

`davinci` **completion**

```
Explain the theory of gravity to a 6 year
old.

Explain the theory of relativity to a 6
year old in a few sentences.

Explain the big bang theory to a 6 year
old.

Explain evolution to a 6 year old.
```

`instruct-davinci` **completion**

```
People went to the moon, and they took
pictures of what they saw, and sent them
back to the earth so we could all see
them.
```

Figure 3: Instruct models[1]

[1] https://beta.openai.com/docs/engines/instruct-series-beta

Tasks

They aggregate 62 text datasets that are publicly available on Tensorflow Datasets. Each dataset is categorized into one of twelve task clusters.

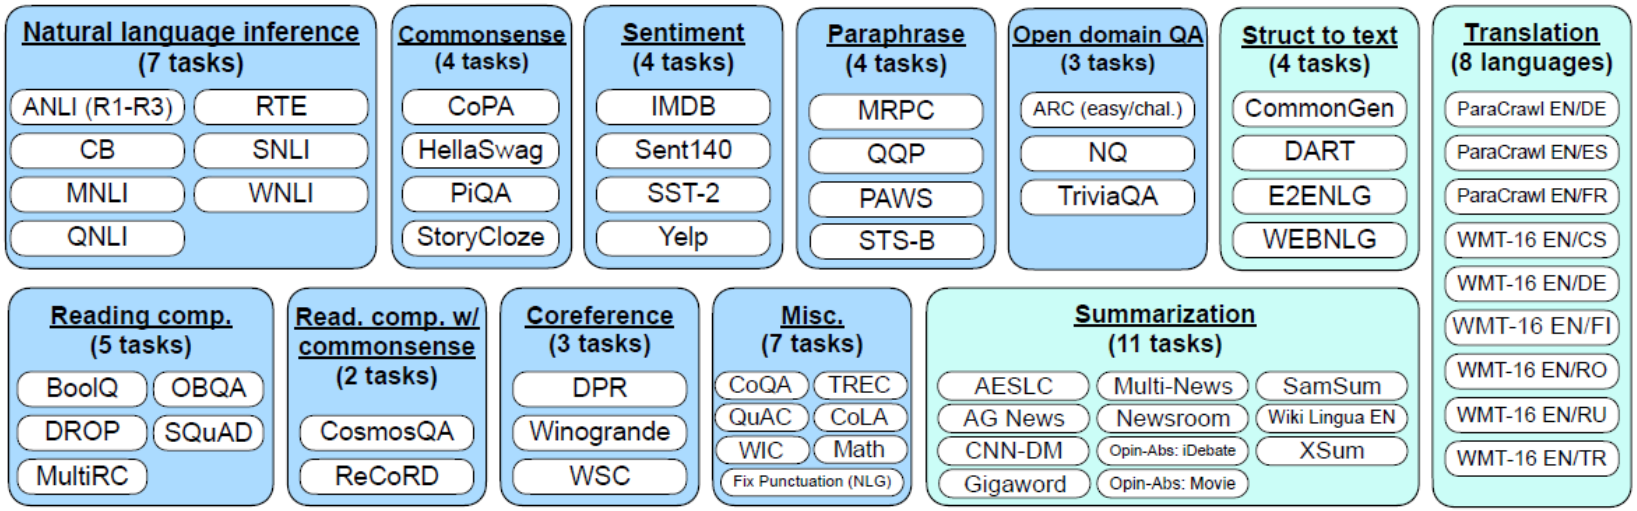

Figure 4: Tasks and clusters

Templates

For each task, they manually compose ten unique templates that describe the task using natural language instructions.

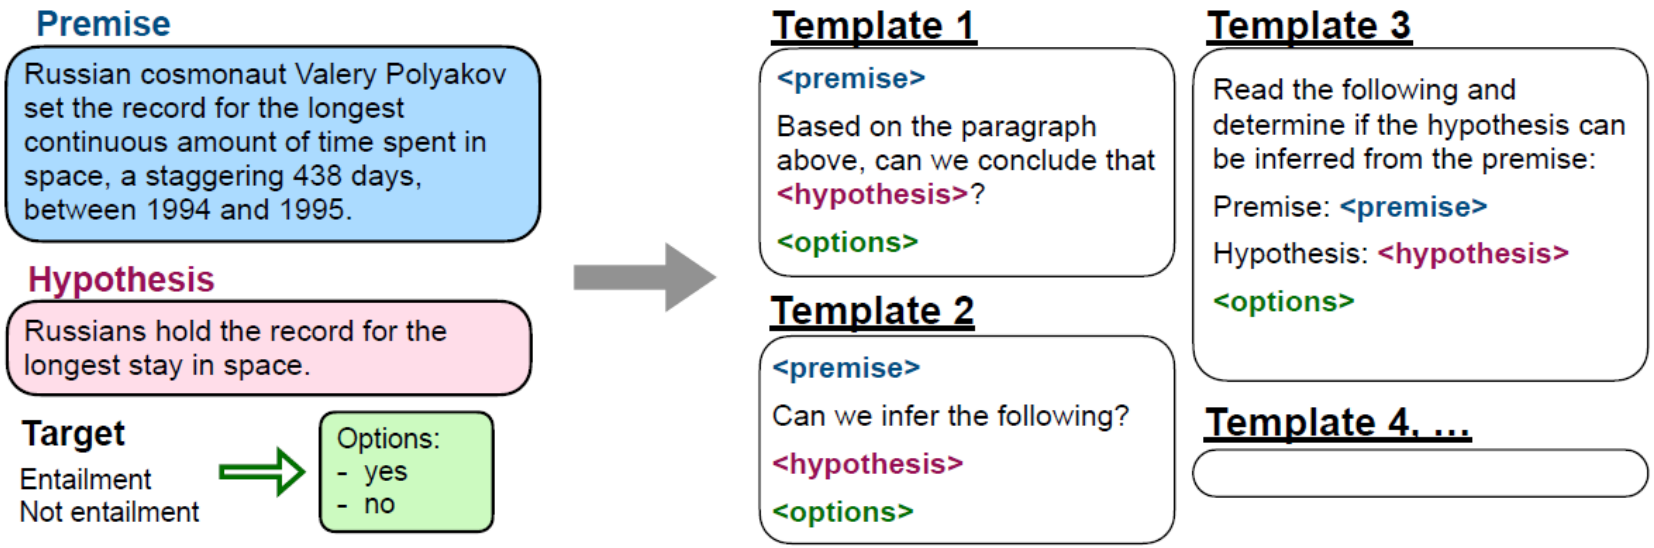

Figure 5: Templates

Input

For classification tasks, they include an options suffix.

INPUT
There are four ways an individual can acquire Canadian citizenship: by birth on Canadian soil; by descent (being born to a Canadian parent); by grant (naturalization); and by adoption. Among them, only citizenship by birth is granted automatically with limited exceptions, while citizenship by descent or adoption is acquired automatically if the specified conditions have been met. Citizenship by grant, on the other hand, must be approved by the Minister of Immigration, Refugees and Citizenship.

Can we conclude that can i get canadian citizenship if my grandfather was canadian?

OPTIONS:
- no
- yes

TARGET
no

Figure 6: Input

Model and Tuning

Model:

- Arch: a dense left-to-right, decoder-only transformer language model of 137B parameters
- Data: 2.81T BPE tokens with a vocabulary of 32K tokens (SentencePiece); Approximately 10% is non-English.

Tuning:

- $>$10m: limit the number to 30000
- low-resource: examples-proportional mixing scheme

Results

	Natural Language Inference				
	ANLI-R1 acc.	ANLI-R2 acc.	ANLI-R3 acc.	CB acc.	RTE acc.
Supervised model	57.4[b]	48.3[b]	43.5[b]	96.8[a]	92.5[a]
Base LM 137B zero-shot	39.6	39.9	39.3	42.9	73.3
· few-shot	39.0	37.5	40.7	34.8	70.8
GPT-3 175B zero-shot	34.6	35.4	34.5	46.4	58.9
· few-shot	36.8	34.0	40.2	82.1	70.4
FLAN 137B zero-shot					
- no prompt engineering	47.7 ▲10.9 stdev=1.4	43.9 ▲8.5 stdev=1.3	47.0 ▲6.8 stdev=1.4	64.1 ↑17.7 stdev=14.7	78.3 ▲7.9 stdev=7.9
- best dev template	46.4 ▲9.6	44.4 ▲9.0	48.5 ▲8.3	83.9 ▲1.8	84.1 ▲13.9

Figure 7: Results on natural language inference.

Cluster

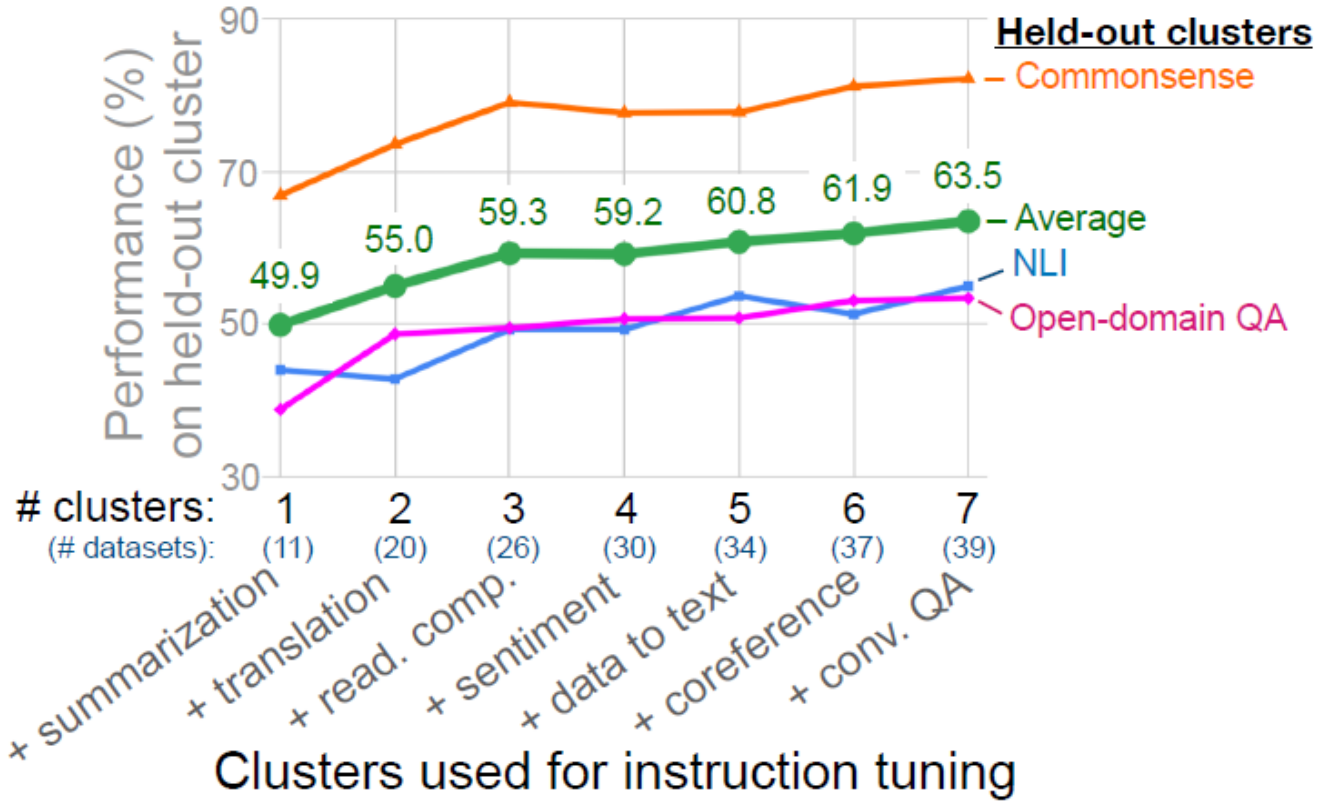

Adding additional task clusters to instruction tuning improves zero-shot performance on held-out task clusters.

Tasks and Templates

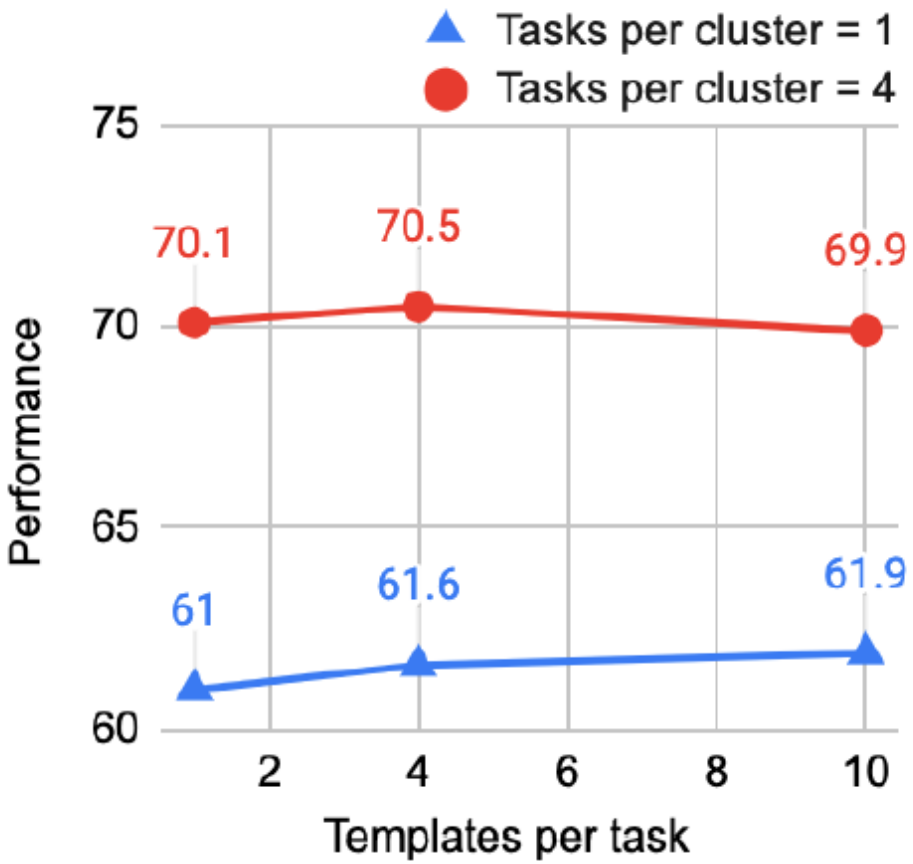

Using more tasks per cluster improved performance.

Using more templates per task, however, had a comparatively negligible effect on performance.

Model Size

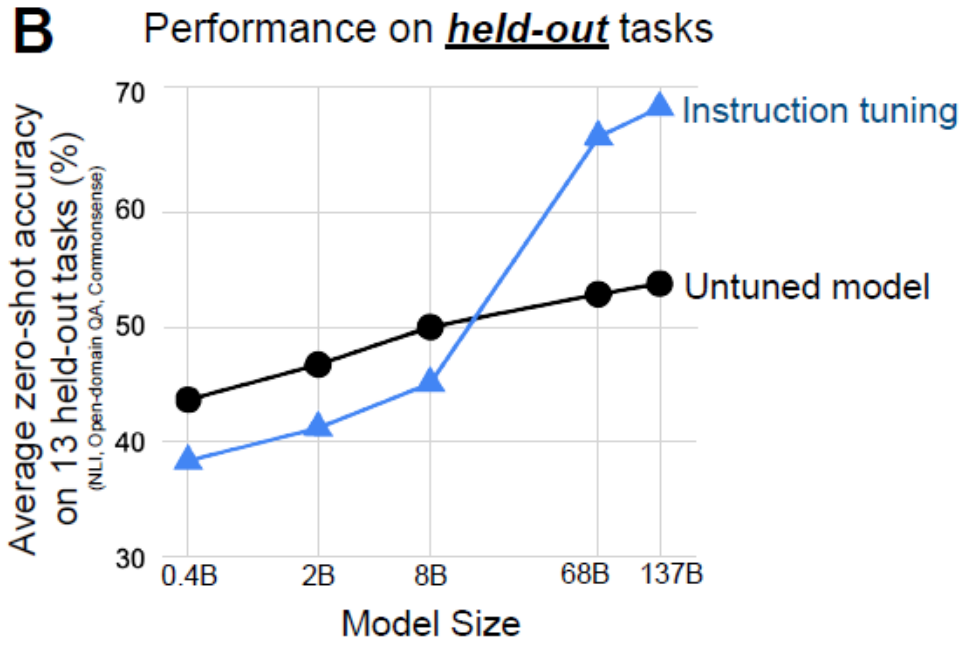

Instruction tuning actually hurts performance on held-out tasks for small-scale models.

Overall

- FLAN substantially improves the performance of its unmodified counterpart and surpasses zero-shot 175B GPT-3 on 19 of 25 tasks.
- FLAN even outperforms few-shot GPT-3 by a large margin on ANLI, RTE, BoolQ, AI2-ARC, OpenbookQA, and StoryCloze.
- Number of tasks and model scale are key components to the success of instruction tuning.

Comparison

- vs T5 and Prompt learning

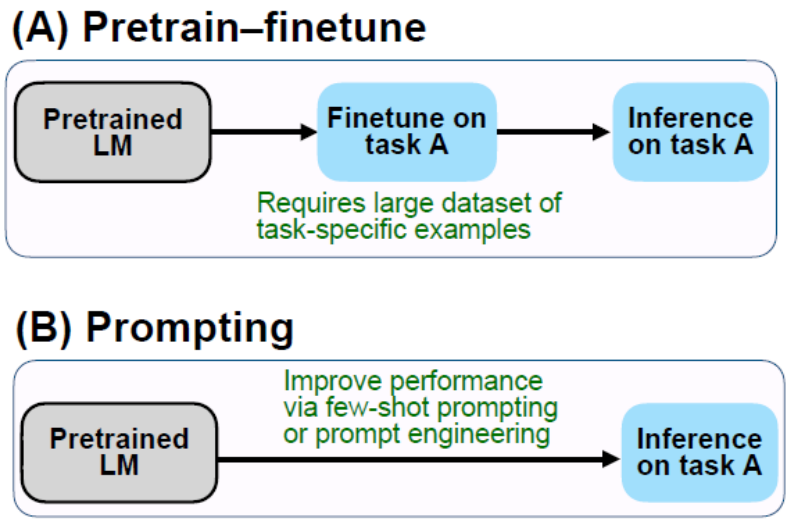

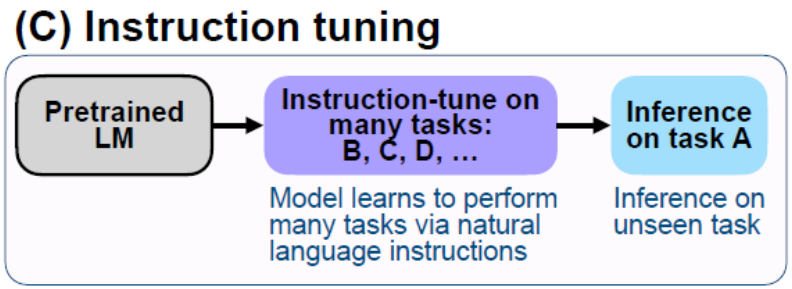

- vs GPT-3 in-context learning

Conclusion

This paper has explored instruction tuning.

- They presented FLAN, a 137B parameter language model that performs NLP tasks described using instructions.
- The performance of FLAN compares favorably against both zero-shot and few-shot GPT-3

Bibliography

[1] J. Wei, M. Bosma, V. Y. Zhao, K. Guu, A. W. Yu, B. Lester, N. Du, A. M. Dai, and Q. V. Le, "Finetuned language models are zero-shot learners," *arXiv preprint arXiv:2109.01652*, 2021.

[2] T. B. Brown, B. Mann, N. Ryder, M. Subbiah, J. Kaplan, P. Dhariwal, A. Neelakantan, P. Shyam, G. Sastry, A. Askell, *et al.*, "Language models are few-shot learners," *arXiv preprint arXiv:2005.14165*, 2020.